

Sai Sathwik Reddy Kondapuram

(470) 437-7322 | sathwikkondapuram@gmail.com | Fort Worth, TX
linkedin.com/in/sathwik-kondapuram | github.com/sathwik-hpe

PROFESSIONAL SUMMARY

Senior Backend and Cloud Developer with 6+ years building enterprise-grade microservices and cloud-native applications. Developed 19+ production microservices in GoLang for HPE GreenLake platform supporting \$1B+ DISA contract managing 900+ services across distributed Kubernetes clusters. Led team of 4 developers architecting scalable solutions. Expert in Kubernetes operators, RESTful APIs, Infrastructure as Code (Terraform, Ansible), and AI/ML integration using LangChain with RAG patterns for intelligent automation. Strong experience in AWS/Azure cloud platforms, Python automation, and Linux system programming.

TECHNICAL SKILLS

Backend Development: GoLang, Python, Node.js, REST APIs, Microservices Architecture, gRPC, API Design

Cloud & Container Platforms: Kubernetes, Docker, AWS (EC2, S3, EKS, Lambda, RDS, EMR), Azure, Helm

Infrastructure as Code: Terraform, Ansible, YAML, Shell Scripting, Automation, Linux Administration

AI/ML Integration: LangChain, RAG (Retrieval Augmented Generation), LLM Tool Integration, Function Calling

Kubernetes Advanced: Custom Operators, CRDs (Custom Resource Definitions), Controllers, Client-go

Databases: PostgreSQL, MongoDB, Redis, Oracle, PL/SQL, Database Design, Query Optimization

Programming: Go, Python, JavaScript, TypeScript, Java, Bash, Shell Scripting

Monitoring: Elasticsearch, Kibana, Prometheus, Grafana, Logging, Performance Monitoring

PROFESSIONAL EXPERIENCE

Hewlett Packard Enterprise (HPE)

July 2022 – Present

Senior Cloud Developer

Houston, TX

- Developed 19+ production microservices in GoLang including backup manager, version manager, certificate management, configuration manager, credentials manager, network settings, health checks, admin control, snapshot manager, support logs collector, download gateway, and webhook service for GreenLake platform supporting \$1B+ DISA contract.
- Led team of 4 developers architecting cloud-native solutions, coordinating with QA, Product Management, and defense/government customers to deliver features managing 900+ microservices.
- Built Kubernetes operators and Custom Resource Definitions (CRDs) in GoLang using client-go library to automate platform lifecycle, enable self-healing, and provide declarative resource management.
- Architected AI-powered operations platform using LangChain with RAG to integrate LLMs with Kubernetes APIs for automated diagnostics and troubleshooting, reducing incident resolution time by 70%.
- Designed LLM integration with function calling and multi-step reasoning to create intelligent automation for complex operational tasks.
- Developed RESTful APIs with OpenAPI/Swagger specifications for authentication, backup/restore, certificate rotation, version management, and cluster administration.
- Built Infrastructure as Code using Terraform and Ansible to automate multi-cloud Kubernetes deployments on AWS and Azure.
- Created Python and Shell automation scripts for Linux system administration and deployment orchestration across distributed environments.
- Implemented centralized logging using Elasticsearch and Kibana to process 20GB+ daily logs with custom indexing for real-time monitoring.
- Designed disaster recovery system achieving 100% data recovery for databases and persistent storage with automated backup verification.
- Developed GitHub Actions CI/CD workflows for automated build, test, containerization, and deployment using multi-stage Docker builds.
- Built host agent in GoLang for system operations and cluster management using gRPC communication.
- Collaborated with customers to gather requirements, present solutions, and deliver zero-downtime production releases.

Tech Stack: GoLang, Python, Kubernetes, Docker, Terraform, Ansible, LangChain, RAG Patterns, AWS, Azure, PostgreSQL, MongoDB, Elasticsearch, Kibana, GitHub Actions, Shell Scripting, Linux, OpenAPI, gRPC

OneTrust	January 2022 – July 2022
<i>Software Engineer</i>	<i>Atlanta, GA</i>
<ul style="list-style-type: none"> Developed responsive UI components using Angular and TypeScript for enterprise privacy and consent management platform with focus on user experience. Built interactive dashboards with Angular services and RxJS observables for reactive data visualization and state management. Designed RESTful APIs using Node.js and Express to handle frontend-backend integration with complex business logic and data workflows. Developed backend services with PostgreSQL for transactional data persistence, query optimization, and relational database modeling. Implemented JWT authentication and OAuth authorization across Angular frontend and Node.js backend. Deployed applications on Azure with automated scaling and monitoring configurations. 	
Tech Stack: <i>Angular, TypeScript, Node.js, Express, RESTful APIs, PostgreSQL, Azure, RxJS, JWT, OAuth</i>	
Amazon	June 2021 – August 2021
<i>Data Engineer Intern</i>	<i>Seattle, WA</i>
<ul style="list-style-type: none"> Designed cloud-native data pipelines using Apache Spark and Apache Hudi on AWS EMR to process large-scale datasets. Built automated ETL workflows with AWS Lambda, S3, and Glue to implement GDPR compliance for customer data retention and privacy. Implemented Infrastructure as Code using CloudFormation templates and CI/CD pipelines with AWS CodePipeline for automated deployments. Developed event-driven architectures using AWS SNS/SQS for asynchronous processing and microservice decoupling. Optimized Apache Hive queries and data partition strategies achieving 90% reduction in compute costs. 	
Tech Stack: <i>Python, Apache Spark, Apache Hudi, AWS (EMR, S3, Lambda, Glue, SNS, SQS), CloudFormation, Apache Hive</i>	
Carl Vinson Institute, University of Georgia	August 2020 – December 2021
<i>Graduate Research Assistant</i>	<i>Athens, GA</i>
<ul style="list-style-type: none"> Performed data analysis using Python with pandas and NumPy to extract insights from large government research datasets. Built interactive dashboards using Qlik Sense for stakeholder reporting and data-driven decision making. Developed backend APIs using Node.js and Python for data integration and ETL pipeline automation. Deployed applications on AWS (EC2, S3) with security configurations and automated deployment scripts. 	
Tech Stack: <i>Python, Node.js, pandas, NumPy, Qlik Sense, AWS (EC2, S3), PostgreSQL, Data Analysis</i>	
Apps Associates	August 2018 – June 2019
<i>Associate Software Engineer</i>	<i>India</i>
<ul style="list-style-type: none"> Received comprehensive training on backend development, AWS cloud services, Linux administration, and CloudStack platform. Designed Oracle database schemas and developed PL/SQL stored procedures, triggers, and functions for enterprise applications. Built RESTful APIs using Java and Node.js to connect Oracle databases with external systems for data integration. Performed database performance tuning with index optimization and query analysis to improve application response times. Developed ETL processes and data migration pipelines with validation and error handling for Oracle systems. 	
Tech Stack: <i>Oracle, PL/SQL, SQL, Java, Node.js, REST APIs, AWS, Linux, CloudStack</i>	
EDUCATION	
University of Georgia	<i>Athens, GA</i>
<i>Master of Science in Computer Science</i>	
<i>Coursework: Distributed Systems, Database Management, Cloud Computing, Software Engineering, Algorithms, Data Structures</i>	